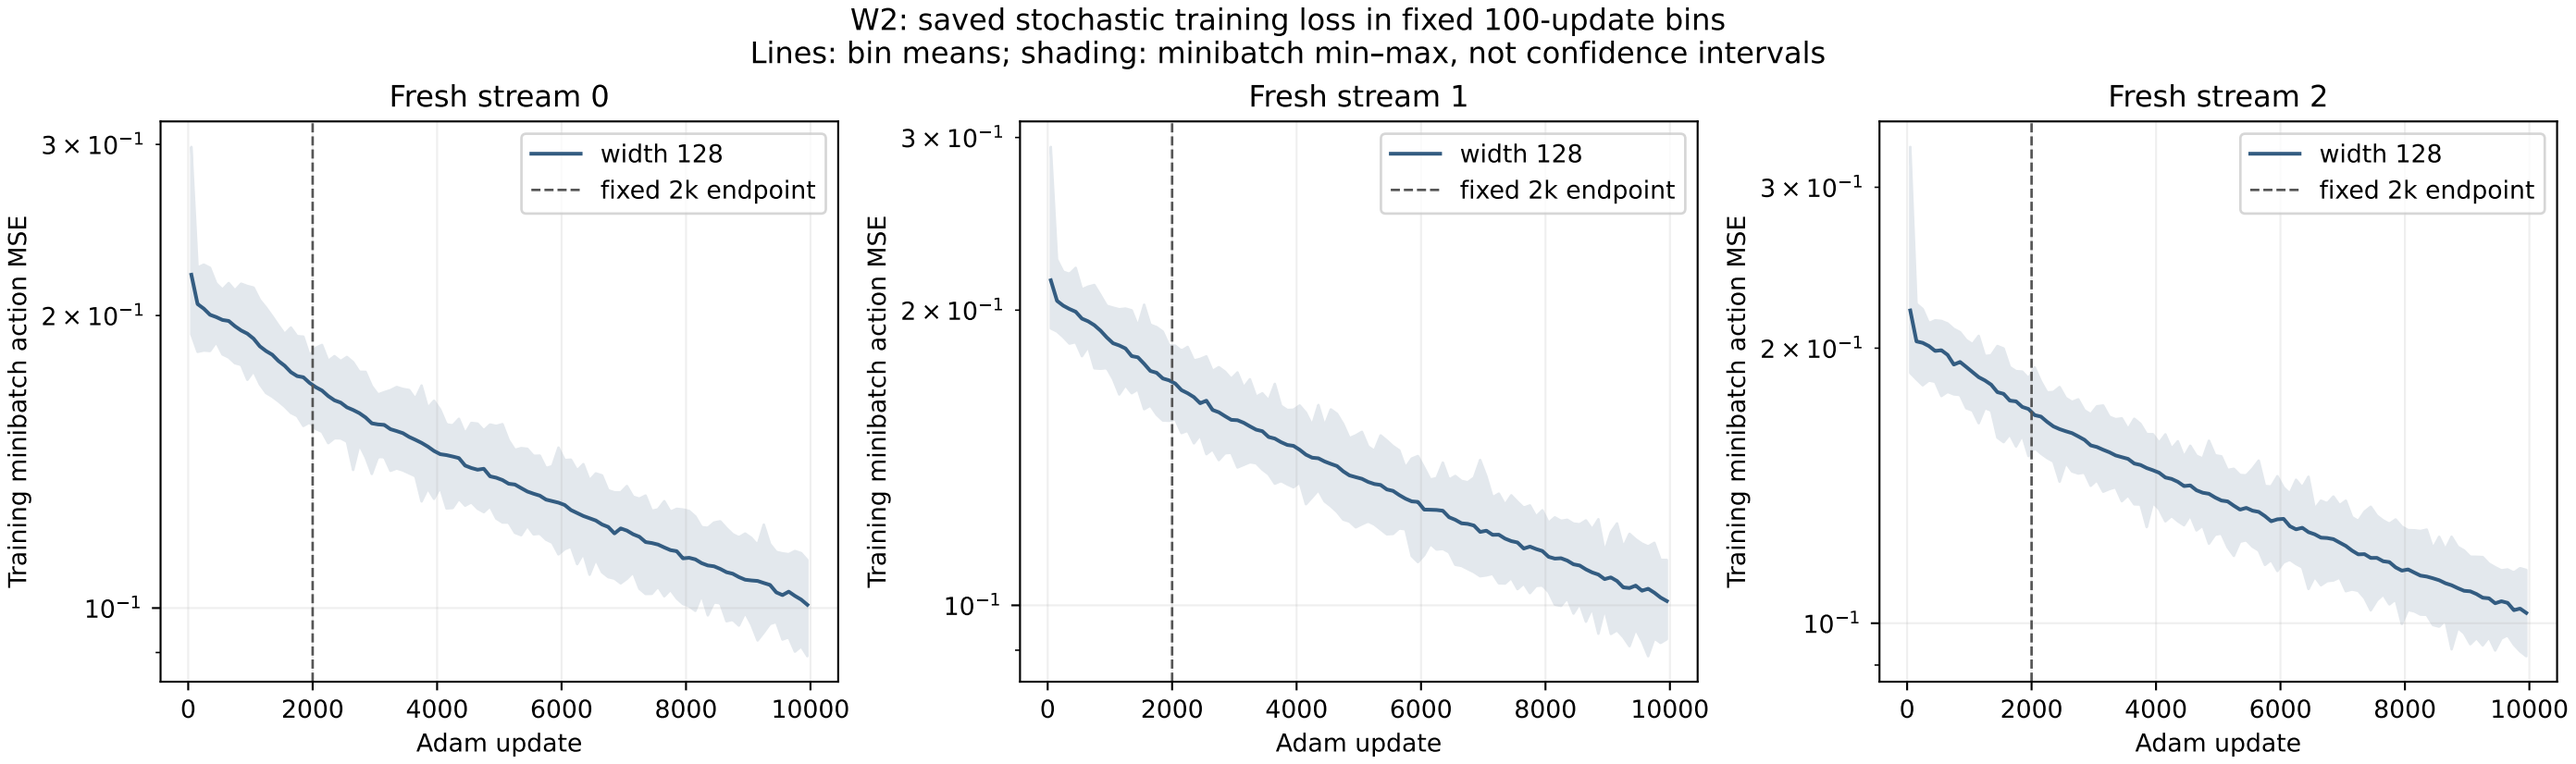
W2: saved stochastic training loss in fixed 100-update bins
Lines: bin means; shading: minibatch min–max, not confidence intervals
Fresh stream 0
Fresh stream 1
Fresh stream 2
width 128
fixed 2k endpoint
Training minibatch action MSE
3×10^{-1}
2×10^{-1}
10^{-1}
0
2000
4000
6000
8000
10000
Adam update